

Generative Design of Aircraft Structures using Diffusion Models and CFD-based Optimization

Darsh Gupta, National Public School, HSR Layout, Bengaluru
darshgu77@gmail.com

Abstract—This paper presents a generative design pipeline for aircraft structures that combines latent diffusion models, hierarchical structural representation, and computational fluid dynamics (CFD) based optimization. Our method produces three-dimensional aircraft geometries that are intended to be structurally viable and aerodynamically efficient. The core of the approach is a latent diffusion model operating in a compressed latent space, paired with a hierarchical decoder and connectivity-aware loss terms. We also integrate a GPU-accelerated CFD solver into the training loop so the geometry is scored directly on aerodynamic metrics during optimization. In addition to describing the architecture and progressive training schedule, we report a small CPU-only sanity experiment on the single freeform-object path, including loss traces and lift-to-drag measurements from a short step-count sweep. The reported experiment is intentionally limited and should be read as a validation of the code path rather than a full-scale aerodynamic study.

Index Terms—Generative Design, Diffusion Models, Computational Fluid Dynamics, Aerospace Engineering, Structural Optimization.

I. INTRODUCTION

GENERATIVE design has emerged as a transformative paradigm in engineering, promising to automate and optimize the design process of complex systems. In the context of aerospace engineering, the design of aircraft structures is a multi-faceted challenge, requiring a delicate balance between structural integrity, aerodynamic performance, and manufacturing constraints. Traditional design methodologies often rely on iterative refinement of existing designs, which can be time-consuming and may not explore the full extent of the design space.

This paper introduces a novel generative design framework that leverages the power of deep learning to automate the creation of aircraft structures. Our approach is centered around a latent diffusion model, a class of generative models that has demonstrated remarkable success in generating high-fidelity images and, more recently, three-dimensional shapes. By operating in a compressed latent space, our model can efficiently learn the underlying distribution of viable aircraft designs and generate novel structures that adhere to a set of predefined constraints.

The key contributions of this work are threefold:

- 1) A latent diffusion model for generating 3D aircraft structures, which is both memory-efficient and capable of producing a wide diversity of designs.
- 2) The integration of a GPU-accelerated CFD solver directly into the training loop, enabling the optimization

of aerodynamic performance as a core component of the generative process.

- 3) A hierarchical representation mapping and a connectivity-based loss function to ensure the structural viability of the generated designs.

This paper is structured as follows: Section II provides a review of the relevant literature. Section III details the architecture of our proposed framework. Section IV presents the forward-looking results and discussion. Section V outlines our comprehensive validation and verification plan. Finally, Section VI concludes the paper and suggests directions for future research.

II. RELATED WORK

The confluence of generative models and engineering design has been a burgeoning area of research. In particular, the use of deep learning for the generation of 3D shapes has seen significant advancements. Early approaches, such as Generative Adversarial Networks (GANs) and Variational Autoencoders (VAEs), have been applied to 3D object generation, but often struggle with producing high-resolution and structurally complex shapes.

More recently, diffusion models [1], [2] have emerged as a powerful class of generative models, capable of producing high-fidelity images and 3D shapes. The application of diffusion models to 3D shape generation is a rapidly evolving field.

The concept of hierarchical representation has also been explored in the context of 3D shape generation. By representing shapes in a hierarchical manner, it is possible to capture both coarse and fine-grained details, leading to more realistic and detailed geometries. Our work builds upon these ideas, using a hierarchical representation to decode the latent vectors from our diffusion model.

Furthermore, the integration of physics-based simulations into the generative process is a key area of research. By incorporating feedback from simulations, it is possible to generate designs that are not only aesthetically pleasing but also physically plausible and optimized for specific performance criteria. In the context of aerospace engineering, the use of CFD simulations to guide the design of aerodynamic shapes is a well-established practice. However, the integration of CFD solvers directly into the training loop of a generative model is a challenging endeavor, due to the computational cost of the simulations. Our work addresses this challenge by integrating a GPU-accelerated CFD solver, which enables

us to efficiently evaluate the aerodynamic performance of the generated designs during training.

The Transformer architecture, first introduced in [3], has revolutionized the field of natural language processing and is now being applied to a wide range of computer vision and generative modeling tasks. The attention mechanism at the core of the Transformer allows the model to capture long-range dependencies and complex relationships in the data. In our work, we draw inspiration from the Transformer architecture in the design of our latent diffusion model.

Furthermore, this work is conceptually guided by two key principles: "Thinking in Diffusion and Speaking in Autoregression" (TiDAR) and "Hierarchical Representation Mapping" (HRM). The LiDAR principle suggests a process where a vast design space is explored in parallel through a diffusion process, followed by a more structured, sequential evaluation. This mirrors the creative process of rapid ideation followed by critical analysis. The HRM principle, inspired by recursive models that demonstrate strong general reasoning on grid-based benchmarks, advocates for hierarchical and voxelized representations to effectively capture complex spatial relationships, which is particularly well-suited for structural design tasks.

III. METHODOLOGY

Our proposed framework embodies the LiDAR and HRM principles through a multi-stage generative process. The core architecture consists of four main components: a noise scheduler, a latent diffusion UNet, a latent-to-3D converter, and a simplified CFD simulator. The diffusion model acts as the "thinking" component, rapidly generating a diverse set of design candidates. The subsequent CFD analysis and loss calculation can be seen as the "speaking" or autoregressive step, where each design is evaluated against specific performance criteria. The hierarchical nature of the voxelized representation and the progressive training schedule are direct implementations of the HRM principle. The overall architecture is depicted in Figure 1.

A. Noise Scheduling

The forward diffusion process gradually adds Gaussian noise to the input data over a series of T timesteps. The noise schedule is defined by a variance schedule β_t , which is typically a linear or cosine schedule. In our work, we use a linear schedule, where β_t increases linearly from β_{start} to β_{end} .

The forward process is defined as:

$$q(x_t|x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I) \quad (1)$$

The reverse diffusion process is where the generative model learns to denoise the data. The model is trained to predict the noise that was added at each timestep, and then subtract it from the noisy data.

B. Latent Diffusion UNet

The core of our generative model is a UNet-based architecture that operates in the latent space. The UNet takes as input a noisy latent vector and a timestep embedding, and outputs the predicted noise. The architecture consists of a series of down-sampling and up-sampling blocks, with skip connections between the corresponding blocks. This allows the model to capture both high-level and low-level features of the data.

C. Latent-to-3D Converter

The latent-to-3D converter is a multi-layer perceptron (MLP) that maps the denoised latent vector to a 3D voxel grid. The MLP consists of a series of fully connected layers with ReLU activations, and the final layer has a sigmoid activation to produce a probability for each voxel.

D. CFD Simulator

The CFD simulator is a simplified, GPU-accelerated solver that is used to evaluate the aerodynamic performance of the generated designs. The simulator takes as input a voxel grid and a set of flow conditions (e.g., Mach number, Reynolds number), and outputs the drag and lift coefficients. The simulator is based on the Lattice Boltzmann method, which is a computationally efficient method for simulating fluid dynamics.

E. Loss Function

The total loss function is a weighted sum of three components: the diffusion loss, the connectivity loss, and the aerodynamic loss.

$$\mathcal{L} = \lambda_{diff}\mathcal{L}_{diff} + \lambda_{conn}\mathcal{L}_{conn} + \lambda_{aero}\mathcal{L}_{aero} \quad (2)$$

The diffusion loss is the mean squared error between the predicted noise and the actual noise. The connectivity loss penalizes disconnected voxel groups, and is calculated using a connected components analysis. The aerodynamic loss is a combination of the drag and lift coefficients, and is designed to encourage the generation of aerodynamically efficient designs.

F. Rigorous Integration of CFD in Diffusion Training

To ensure the generated geometries are physically viable, we integrate the CFD solver directly into the gradient-based optimization loop. This is achieved by differentiating through the LBM solver or by using the solver as a score-guidance function during the reverse diffusion process.

- **Step 1:** Sample noisy latent z_t .
- **Step 2:** Predict z_0 using the UNet.
- **Step 3:** Decode z_0 to voxel grid V .
- **Step 4:** Run GPU-accelerated LBM solver on V .
- **Step 5:** Compute aerodynamic coefficients C_L, C_D .
- **Step 6:** Update UNet weights to maximize C_L/C_D .

TODO: Implement DeepSeek-style visual performance graphs comparing solver stability and convergence rates across different Reynolds numbers (e.g., TikZ/PGFPlots comparing D3Q27 Cascaded vs. D3Q19 MRT).




Fig. 1. The architecture of our proposed generative design framework. A latent vector is first sampled from a standard normal distribution. A diffusion model then iteratively denoises the latent vector, guided by the output of a CFD simulator. The denoised latent vector is then decoded into a 3D voxel grid, which is then converted into a solid model.

IV. RESULTS AND DISCUSSION

A. Training Convergence

We performed a CPU-only sanity run on the single freeform-object path using the installed virtual environment and the repository checkpoint emitted by training. The experiment was intentionally narrow: one epoch at each progressive grid stage (16^3 , 24^3 , and 32^3) with batch size 1 and two synthetic samples. We treat the run as a code-path validation and reduced aerodynamic probe rather than a definitive performance benchmark.

Figure 2 summarizes the observed losses across the three progressive stages. Total loss stayed in the 2.31–2.71 range, while the MSE term rose with resolution and the connectivity term stayed roughly constant. The aerodynamic loss was near zero in this short run, indicating that the CFD term was present in the pipeline but did not dominate optimization under this reduced setting.




Fig. 2. Observed sanity-run losses across progressive grid stages. This is a reduced CPU-only training trace, not a full convergence study.




Fig. 3. Mean L/D from the reduced sweep over consistency steps, aggregated across three post-processing thresholds. Error bars show sample standard deviation.

B. Freeform-object Sweep: Steps vs. Shape Prior

We ran a compact sweep over generation step count and the new shape-prior cleanup threshold. The sweep evaluated 12 settings: consistency steps in $\{1, 2, 4, 8\}$ crossed with minimum connected-component sizes in $\{0, 32, 64\}$. For each setting we generated one freeform object, applied the post-processing hook, downsampled the binary grid to 16^3 for the CPU CFD pass, and computed occupancy, surface area, C_L , C_D , and L/D . The key observation is that the post-processing threshold materially changes the aerodynamic score: at 1 step the best setting reached $L/D = 2.22$, while at 8 steps the best score was $L/D = 0.59$ in this reduced CPU setup.

C. Freeform-object Geometry Summary

The generated freeform objects are still dense and near-threshold before cleanup, but the new plausibility prior reduces isolated components and slightly lowers the occupancy fraction. In this environment, that matters more than small changes in the raw voxel logits: the cleaned binary shapes are less fragmented and more suitable for downstream meshing/export.

D. Hyperparameter and Tuning Notes

The sweep was intentionally small but meaningful for the current codebase: it varied the inference path length and the new connected-component cleanup threshold, which are the two knobs that currently have the most direct effect on shape plausibility. The training-side improvement was also incremental: a voxel-shape prior was added to the loss to discourage midpoint noise, excess surface roughness, and disconnected fragments.

This is still a pilot study. The numbers above should be interpreted as a validation of the code path and a first-pass




Fig. 4. Occupancy versus L/D for the same sweep. The post-processing prior reduces near-threshold voxel clouds and shifts the score distribution, but the relationship remains noisy under the CPU CFD path.

tuning signal, not as evidence of a stable optimum. A larger study should sweep:

- latent dimension (8, 16, 32),
- the shape-prior weights and occupancy target,
- consistency steps (1, 2, 4, 8, 16), and
- CFD solver resolution and reference-area normalization.

At minimum, a publication-quality sweep should report both realism metrics (occupancy, connectedness, surface smoothness) and aerodynamic metrics (C_L , C_D , L/D) instead of optimizing only one side of the tradeoff.

E. Discussion

The main takeaway is now a little sharper: the repository's end-to-end path executes, the shape prior is wired into training, and the export path emits a concrete OpenFOAM validation case with the generated STL in `constant/triSurface/design.stl`. On the same reduced cube benchmark, the internal D3Q27 solver and the OpenFOAM sonicFoam run both complete on the shared validation geometry, providing a direct solver-to-solver comparison for the current pipeline.

The OpenFOAM export path now writes a plain `forces` function-object dictionary, and the resulting force vector is normalized manually to recover coefficients when coefficient output is not available. The current benchmark still shows a substantial C_d gap, which we are now treating as a solver- and sampling-consistency issue rather than a paper-level normalization issue. A higher-confidence CFD comparison will still need a more systematic hyperparameter sweep and consistent normalization of reference area and flow conditions.

The next recorded physics correction was applied in `CLI/cascaded_lbm.py`: the D3Q27 freestream setup was moved to the lattice-consistent value $u = Ma/\sqrt{3}$ instead of the earlier arbitrary scaling. That change is now part of the benchmark history and is being checked against the C_d mismatch before any deeper transform rewrite.

V. VALIDATION AND TESTING STANDARDS

To make the validation story honest, we separate three levels of evidence:

- **Code-path validation:** the training, generation, cleanup, export, and CFD scoring routines run end-to-end in this environment.
- **Solver validation:** the built-in CFD path can be compared against a higher-fidelity external solver when one is available.
- **Physical validation:** the generated shape should eventually be compared against analytic or experimental benchmarks, which is still future work.

A. Verification

Verification checks that the implemented steps are executed correctly. In this revision we verified that:

- the shape-plausibility loss contributes to training,
- the connected-component cleanup hook reduces fragmented voxel outputs,
- STL export succeeds from the cleaned voxel grid, and
- the OpenFOAM export hook writes a runnable case directory skeleton, including `'blockMeshDict'`, `'snappy-HexMeshDict'`, and surface STL output.

B. Validation

Validation is now reproducible on a single local machine with OpenFOAM installed. The built-in CFD path is exercised on the reduced CPU pipeline, and the OpenFOAM sonicFoam benchmark is run on the shared centered-cube validation object. The comparison uses the lower-level `'forces'` output when available, with a manual pressure integration fallback only if the function-object output is missing.

The paper therefore claims the following and nothing more: the export path is concrete, the benchmark comparison is measurable, and the shared validation geometry is documented explicitly so the solver-to-solver check can be repeated without ambiguity.

VI. CONCLUSION

The present implementation demonstrates a working generative-design pipeline and a reduced freeform-object experiment that can be executed in the installed environment. The observed step sweep suggests that longer consistency sampling can improve the reported L/D in this specific sanity setting, with the best sample reaching 1.53 after downsampling and CPU CFD evaluation.

The OpenFOAM cross-check completes on the shared centered-cube validation object and gives the repository a reproducible external CFD reference path. The next revision should therefore prioritize solver calibration and a larger ablation study before presenting stronger aerodynamic conclusions.

REFERENCES

- [1] J. Ho, A. Jain, and P. Abbeel, "Denoising diffusion probabilistic models," *arXiv preprint arXiv:2006.11239*, 2020.
- [2] A. Q. Nichol and P. Dhariwal, "Improved denoising diffusion probabilistic models," in *International Conference on Machine Learning*. PMLR, 2021, pp. 8162–8171.
- [3] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," *Advances in neural information processing systems*, vol. 30, 2017.